W4111 – Introduction to Databases
Section 003, V03, Fall 2024
Lecture 4: ER(3), Relational(3), SQL(3)

Columbia | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

W4111 – Introduction to Databases
Section 003, V03, Fall 2024
Lecture 4: ER(3), Relational(3), SQL(3)

We will start in a few minutes.

Columbia ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Today's Contents

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Contents

- Introduction and course updates.

- ER Modeling (3).

- Relational model and algebra (3).

- SQL (3).

- Homework and projects.

Columbia Engineering
The Fu Foundation School of Engineering and Applied Science

Introduction
and
course updates

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Course Updates

ER Modeling

© Donald F. Ferguson, 2024

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Let's Examine and Do Some Examples

Columbia | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Relational Model

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Equivalent Queries

- There is more than one way to write a query in relational algebra.

- Example: Find information about courses taught by instructors in the Physics department with salary greater than 90,000

- Query 1

$$\sigma_{\text{dept\_name="Physics"} \wedge \text{salary} > 90,000} (instructor)$$

- Query 2

$$\sigma_{\text{dept\_name="Physics"}} (\sigma_{\text{salary} > 90.000} (instructor))$$

- The two queries are not identical; they are, however, equivalent -- they give the same result on any database.

Equivalent Queries

- There is more than one way to write a query in relational algebra.

- Example: Find information about courses taught by instructors in the Physics department

- Query 1

  $\sigma_{dept\_name=\text{"Physics"}} (instructor \bowtie_{instructor.ID \, = \, teaches.ID} teaches)$

- Query 2

  $(\sigma_{dept\_name=\text{"Physics"}} (instructor)) \bowtie_{instructor.ID \, = \, teaches.ID} teaches$

- The two queries are not identical; they are, however, equivalent -- they give the same result on any database.

What are all those other Symbols?

- τ   order by
- γ   group by
- ¬   negation
- ÷   set division
- ⋈   natural join, theta-join
- ⟕   left outer join
- ⟖   right outer join
- ⟗   full outer join
- ⋉   left semi join
- ⋊   right semi join
- ▷   anti-join

- Some of the operators are useful and "common," but not always considered part of the core algebra.
- Some of these are pretty obscure
  - Division
  - Anti-Join
  - Left semi-join
  - Right semi-join
- Most SQL engines do not support them.
  - You can implement them using combinations of JOIN, SELECT, WHERE, … …
  - But, I cannot every remember using them in applications I have developed.
- Outer JOIN is very useful, but less common. We will cover.
- There are also some "patterns" or "terms"
  - Equijoin
  - Non-equi join
  - Natural join
  - Theta join
  - … …
- I may ask you to define these terms on some exams or the obscure operators because they may be common internships/job interview questions.
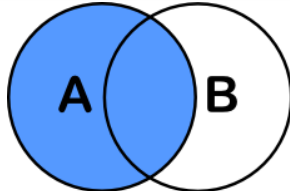
Columbia | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

What are all those other Symbols?

- τ   order by
- γ   group by
- ¬   negation
- ÷   set division
- ⋈   natural join, theta-join
- ⟕   left outer join
- ⟖   right outer join
- ⟗   full outer join
- ⋉   left semi join
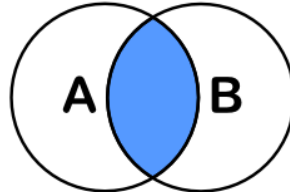- ⋊   right semi join
- ▷   anti-join

```
1  σ student_dept='Comp. Sci.' ∨ student_dept='Elec. Eng.'
2  (
3    π student_id, student_name, student_dept,
4      advisor_id, advisor_name, advisor_dept
5    (
6      (
7        π student_id←ID, student_name←name,
8          student_dept←dept_name, i_id
9            (student ⋈ ID=s_id advisor)
10     )
11     ⋈ i_id=advisor_id
12     (
13       π advisor_id←ID, advisor_name←name, advisor_dept←dept_name
14           (instructor ⋈ ID=i_id advisor)
15     )
16   )
17 )
```

student_id	student_name	student_dept	advisor_id	advisor_name	advisor_dept
128	Zhang'	Comp. Sci.'	45565	Katz'	Comp. Sci.'
12345	Shankar'	Comp. Sci.'	10101	Srinivasan'	Comp. Sci.'
76543	Brown'	Comp. Sci.'	45565	Katz'	Comp. Sci.'
76653	Aoi'	Elec. Eng.'	98345	Kim'	Elec. Eng.'
98765	Bourikas'	Elec. Eng.'	98345	Kim'	Elec. Eng.'

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

One Way to Think About Joins

SELECT <auswahl>
FROM tabelleA A
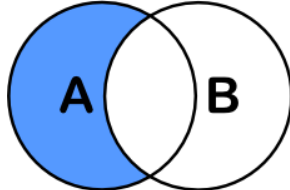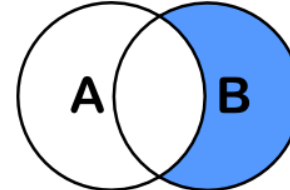LEFT JOIN tabelleB B
ON A.key = B.key

CHEATSHEET SQL JOINS

SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key
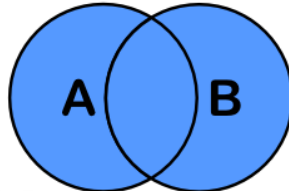
SELECT <auswahl>
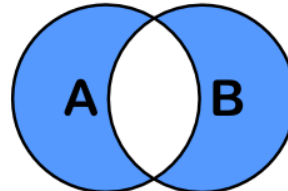FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
WHERE B.key IS NULL

SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL

SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
OR B.key IS NULL

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science
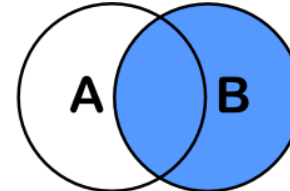
Thinking about JOINs

- Some terms:
  - Natural Join
    - Equality of A and B columns
    - With the same name.
  - Equijoin
    - Explicitly specify columns that must have the same value.
    - A.x=B.z AND A.q=B.w
  - Theta Join: Arbitrary predicate.
- Inner Join
  - JOIN "matches" rows in A and B.
  - Result contains ONLY the matched pairs.
- What I want is:
  - All the rows that matched.
  - And the rows from A that did not match?
  - OUTER JOIN (⋈, ⋈)

SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
WHERE B.key IS NULL

SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL

SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
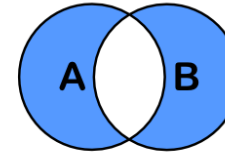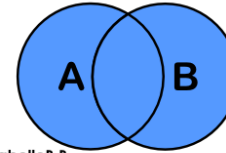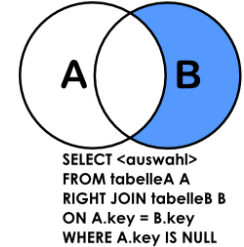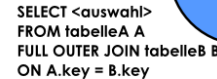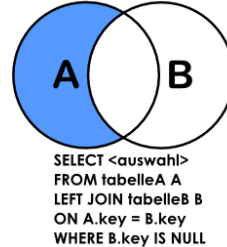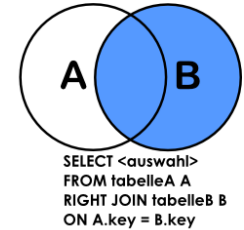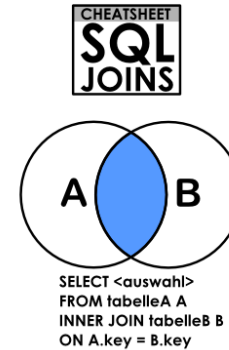OR B.key IS NULL

Columbia | Engineering
The Fu Foundation School of Engineering and Applied Science

Do Some Joins

- prereq is interesting
  - All courses with prereqs.
  - All courses and prereqs even if it does not have a prereq.
  - All courses and ones that are prereqs
  - Courses that do not have prereqs
  - Etc.
- See notebook with examples.

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Some Examples

- **A course and its prerequisites**

  π course_id←course.course_id,

  course_title←course.title,

  prerequisite_id←prereq.prereq_id,

  prerequisite_title←p.course_id

  (

  (course ⋈ prereq)

  ⋈ prereq.prereq_id=p.course_id

  ρ p(course))

- Join course and prereq to get
  - course info
  - prereq_id
- Join with
  - course using prereq_id
  - To get the prereq info
- But,
  - course.course_id is ambiguous
  - Because the table appears twice.
  - So, I have to "alias" the second use.
- Also, note the use of renaming in the project for column names.

Columbia | Engineering
The Fu Foundation School of Engineering and Applied Science

- **Courses-prerequisites and courses without prereqs.**

  τ prerequisite_id

  (π course_id←course.course_id,

        course_title←course.title,

        prerequisite_id←prereq.prereq_id,

        prerequisite_title←p.course_id

  (

        (course ⋈ prereq)

        ⋈ prereq.prereq_id=p.course_id

        ρ p(course))

)

Some Examples

- **Courses that are not prerequisites**

  τ prereq_course_id

  (

    π prereq_course_id←prereq.course_id,

      course_id←course.course_id,

      course_title←course.title

      (

        prereq ⋈ prereq_id=course.course_id course

)

)

Columbia | Engineering
The Fu Foundation School of Engineering and Applied Science

Relational Algebra

- We will do more examples in lectures and HW assignments.

- The language *is interesting but* can be tedious and confusing.

- Useful in some advanced scenarios:
  - Designing query languages for new databases and data models.
  - Understanding how DBMS implement query processing and optimization.

- Also cover because it does come up in some advanced courses, and may come up in job interviews.

- You will be able to use a "cheat sheet" on exams. We will focus on "did you know what to do" and not "did you get the syntax completely correct."

Codd's 12 Rules
NULL

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Codd's 12 Rules

Rule 1: Information Rule

The data stored in a database, may it be user data or metadata, must be a value of some table cell. Everything in a database must be stored in a table format.

Rule 2: Guaranteed Access Rule

Every single data element (value) is guaranteed to be accessible logically with a combination of table-name, primary-key (row value), and attribute-name (column value). No other means, such as pointers, can be used to access data.

Rule 3: Systematic Treatment of NULL Values

The NULL values in a database must be given a systematic and uniform treatment. This is a very important rule because a NULL can be interpreted as one the following – data is missing, data is not known, or data is not applicable.

Rule 4: Active Online Catalog

The structure description of the entire database must be stored in an online catalog, known as data dictionary, which can be accessed by authorized users. Users can use the same query language to access the catalog which they use to access the database itself.

Rule 5: Comprehensive Data Sub-Language Rule

A database can only be accessed using a language having linear syntax that supports data definition, data manipulation, and transaction management operations. This language can be used directly or by means of some application. If the database allows access to data without any help of this language, then it is considered as a violation.

Rule 6: View Updating Rule

All the views of a database, which can theoretically be updated, must also be updatable by the system.

Columbia | Engineering
The Fu Foundation School of Engineering and Applied Science

Codd's 12 Rules

Rule 7: High-Level Insert, Update, and Delete Rule

A database must support high-level insertion, updation, and deletion. This must not be limited to a single row, that is, it must also support union, intersection and minus operations to yield sets of data records.

Rule 8: Physical Data Independence

The data stored in a database must be independent of the applications that access the database. Any change in the physical structure of a database must not have any impact on how the data is being accessed by external applications.

Rule 9: Logical Data Independence

The logical data in a database must be independent of its user's view (application). Any change in logical data must not affect the applications using it. For example, if two tables are merged or one is split into two different tables, there should be no impact or change on the user application. This is one of the most difficult rule to apply.

Rule 10: Integrity Independence

A database must be independent of the application that uses it. All its integrity constraints can be independently modified without the need of any change in the application. This rule makes a database independent of the front-end application and its interface.

Rule 11: Distribution Independence

The end-user must not be able to see that the data is distributed over various locations. Users should always get the impression that the data is located at one site only. This rule has been regarded as the foundation of distributed database systems.

Rule 12: Non-Subversion Rule

If a system has an interface that provides access to low-level records, then the interface must not be able to subvert the system and bypass security and integrity constraints.

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Codd's 12 Rules

Rule 3: Systematic treatment of null values:

- "Null values (distinct from the empty character string or a string of blank characters and distinct from zero or any other number) are supported in fully relational DBMS for representing **missing information** and **inapplicable** information in a systematic way, independent of data type."

- Sometimes programmers and database designers are tempted to use "special values" to indicate unknow, missing or inapplicable values.
  - String: "", "NA", "UNKNOWN", ...
  - Numbers: -1, 0, -9999

- Indicators can cause confusion because you have to carefully code some SQL statements to the specific, varying choices programmers made.

Columbia ENGINEERING
The Fu Foundation School of Engineering and Applied Science

NULL and Correct Answers

```
In [4]:    1  %sql describe aaaaS21Examples.null_examples;
```

```
* mysql+pymysql://dbuser:***@localhost
3 rows affected.
```

Out[4]:

Field	Type	Null	Key	Default	Extra
name	varchar(32)	NO	PRI	None	
weight	int	YES		None	
net_worth	int	YES		None	

```
In [5]:    1  %sql select * from aaaaS21Examples.null_examples;
```

```
* mysql+pymysql://dbuser:***@localhost
4 rows affected.
```

Out[5]:

name	weight	net_worth
Joe	100	100
Larry	0	0
Pete	None	None
Tim	200	200

Without NULL, to get a correct answer:

- I must understand the domain to determine "unknown" values or know what choice a developer made.

- Explicitly include "where weight != 0" in all statements.

- And this varies from column to column, table to table, schema to schema, etc.

```
In [7]:    1  %%sql select avg(weight) as avg_weight, avg(net_worth) as avg_net_worth
           2      from aaaaS21Examples.null_examples where name in ('Joe', 'Larry', 'Tim')
```

```
* mysql+pymysql://dbuser:***@localhost
1 rows affected.
```

Out[7]:

avg_weight	avg_net_worth
100.0000	100.0000

```
In [9]:    1  %%sql select avg(weight) as avg_weight, avg(net_worth) as avg_net_worth
           2      from aaaaS21Examples.null_examples where name in ('Joe', 'Pete', 'Tim')
```

```
* mysql+pymysql://dbuser:***@localhost
1 rows affected.
```

Out[9]:

avg_weight	avg_net_worth
150.0000	150.0000

Columbia | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Null Values

- It is possible for tuples to have a null value, denoted by **null**, for some of their attributes

- **null** signifies an unknown value or that a value does not exist.

- The result of any arithmetic expression involving **null** is **null**
  - Example: 5 + **null** returns **null**

- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

    select *name*
    from *instructor*
    where *salary* **is null**

- The predicate **is not null** succeeds if the value on which it is applied is not null.

Note:
- NULL is an extremely important concept.
- You will find it hard to understand for a while.

Note that RelaX uses "=" for NULL despite the fact that SQL does not.

Null Values (Cont.)

- SQL treats as **unknown** the result of any comparison involving a null value (other than predicates **is null** and **is not null**).
  - Example*: 5 < **null**   or   **null** <> **null**   or   **null** = **null**

- The predicate in a **where** clause can involve Boolean operations (**and**, **or**, **not**); thus the definitions of the Boolean operations need to be extended to deal with the value **unknown**.
  - **and** : *(true* **and** *unknown) = unknown,*
    (false **and** *unknown) = false,*
    (unknown **and** *unknown) = unknown*
  - **or:**   (*unknown* **or** *true*)   = *true*,
    (*unknown* **or** *false*) = *unknown*
    (*unknown* **or** *unknown) = unknown*

- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

SQL

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Aggregate Functions

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

    avg: average value
    min: minimum value
    max: maximum value
    sum: sum of values
    count: number of values

Note: Some database implementations have additional aggregate functions.

Aggregate Functions – Group By

- Find the average salary of instructors in each department
  - **select** *dept_name*, **avg** (*salary*) **as** *avg_salary*
    from *instructor*
    group by *dept_name*;

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000

Another View

Employees

DEPARTMENT_ID	SALARY
10	5500
20	15000
20	7000
30	12000
30	5100
30	4900
30	5800
30	5600
40	7500
40	8000
50	9000
50	8500
50	9500
50	8500
50	10500
50	10000
50	9500

5500
22000

33400 *Sum of Salary in Employees table for each department*

15500

65550

DEPARTMENT_ID	SUM(SALARY)
10	5500
20	22000
30	33400
40	15500
50	65550

- GROUP BY column list
  - Forms partitions containing multiple rows.
  - All rows in a partition have the same values for the GROUP BY columns.
- The aggregate functions
  - Merge the non-group by attributes, which may differ from row to row.
  - Into a single value for each attribute.
- The result is one row per distinct set of GROUP BY values.
- There may be multiple non-GROUP BY COLUMNS, each with its own aggregate function.
- You can use HAVING in place of WHERE on the GROUP BY result.

Columbia | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Aggregate Functions Examples

- Find the average salary of instructors in the Computer Science department

  - **select avg** (*salary*)
    from *instructor*
    where *dept_name*= 'Comp. Sci.';

- Find the total number of instructors who teach a course in the Spring 2018 semester

  - **select count** (**distinct** *ID*)
    from *teaches*
    where *semester* = 'Spring' **and** *year* = 2018;

- Find the number of tuples in the *course* relation

  - **select count** (*)
    from *course*;

Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list

  - /* erroneous query */
    select *dept_name*, *ID*, **avg** (*salary*)
    from *instructor*
    group by *dept_name*;

Aggregate Functions – Having Clause

- Find the names and average salaries of all departments whose average salary is greater than 42000

    > **select** *dept_name*, **avg** (*salary*) **as** *avg_salary*
    > **from** *instructor*
    > **group by** *dept_name*
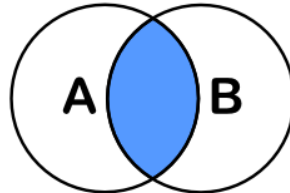   > **having avg** (*salary*) > 42000;

- Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

JOINs
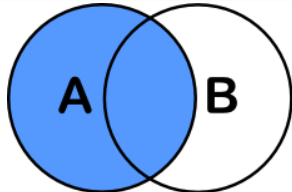
COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

One Way to Think About Joins

SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key

CHEATSHEET
SQL JOINS

SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
WHERE B.key IS NULL

SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key

SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
OR B.key IS NULL

SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL

Columbia | Engineering
The Fu Foundation School of Engineering and Applied Science
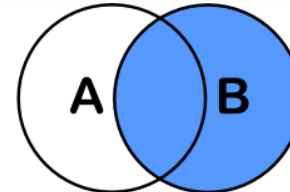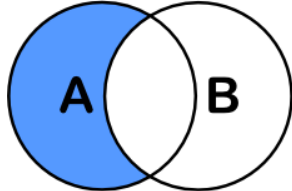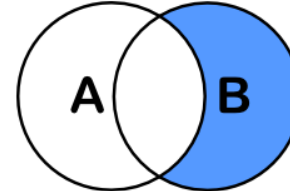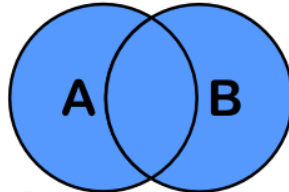
Outer Join

- An extension of the join operation that avoids loss of information.

- Computes the join and then adds tuples form one relation that does not match tuples in the other relation to the result of the join.

- Uses *null* values.

- Three forms of outer join:
  - left outer join
  - right outer join
  - full outer join

Outer Join Examples

- Relation *course*

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

- Relation *prereq*

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that

  course information is missing CS-437

  prereq information is missing CS-315

- *x*

Left Outer Join

- *course* **natural left outer join** *prereq*

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	*null*

- In relational algebra: *course* ⟕ *prereq*

Right Outer Join

- *course* **natural right outer join** *prereq*

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

- In relational algebra: *course* ⟖ *prereq*

Full Outer Join

- *course* **natural full outer join** *prereq*

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	*null*
CS-347	*null*	*null*	*null*	CS-101

- In relational algebra: *course* ⟗ *prereq*

Joined Types and Conditions

- **Join operations** take two relations and return as a result another relation.

- These additional operations are typically used as subquery expressions in the **from** clause

- **Join condition** – defines which tuples in the two relations match.

- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

Join types
inner join
left outer join
right outer join
full outer join

Join conditions
natural
on $<$ predicate $>$
using $(A_1, A_2, \ldots, A_n)$

Joined Relations – Examples

- *course* **natural right outer join** *prereq*

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

- *course* **full outer join** *prereq* **using** (*course_id*)

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101

Joined Relations – Examples

- *course* **inner join** *prereq* **on**
  course.course_id = prereq.course_id

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- What is the difference between the above, and a natural join?

- *course* **left outer join** *prereq* **on**
  course.course_id = prereq.course_id

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	null	null

Joined Relations – Examples

- *course* **natural right outer join** *prereq*

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

- *course* **full outer join** *prereq* **using** (*course_id*)

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101

Let's Do Some Examples

- JOINs using Classic Models

- See ./join-and-aggregate.sql

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Set Operations

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Set Operations

- Find courses that ran in Fall 2017 or in Spring 2018

  (**select** *course_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2017)
  union
  (**select** *course_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2018)

- Find courses that ran in Fall 2017 and in Spring 2018

  (**select** *course_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2017)
  intersect
  (**select** *course_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2018)

- Find courses that ran in Fall 2017 but not in Spring 2018

  (**select** *course_id* **from** *section* **where** *sem* = 'Fall' **and** *year* = 2017)
  except
  (**select** *course_id* **from** *section* **where** *sem* = 'Spring' **and** *year* = 2018)

Note:
- Show in Relax calculator.
- Not all databases support INTERSECT and EXCEPT
- We will see how to handle below.

Set Operations (Cont.)

- Set operations **union, intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the
  - **union all**,
  - **intersect all**
  - **except all**.

Subqueries
Concepts and Examples
(Including Set Membership)

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.

- The nesting can be done in the following SQL query

$$\textbf{select } A_1, A_2, ..., A_n$$
$$\textbf{from } r_1, r_2, ..., r_m$$
$$\textbf{where } P$$

as follows:

- **From clause:** r_i can be replaced by any valid subquery

- **Where clause:** P can be replaced with an expression of the form:

    B <operation> (subquery)

    B is an attribute and <operation> to be defined later.

- **Select clause:**

    A_i can be replaced be a subquery that generates a single value.

Note:

- This is a little cryptic.
- I think I know what they mean.
- There are some operations we will see later in the material, e.g IN, EXISTS, ...

Nested Subquery

- The slides that come with the book have surprisingly little material on nested subqueries.

- The concept is:
  - Extremely important.
  - Students often find subqueries more confusing than joins.
  - The relationship/difference of subqueries to joins is often, initial unclear.

- We have seen:
  - Union sort of puts a table on top of a table.
  - Join puts tables sort of puts tables side-by-side.
  - Subquery enables one query to call another during execution like a subfunction.

Columbia | Engineering
The Fu Foundation School of Engineering and Applied Science
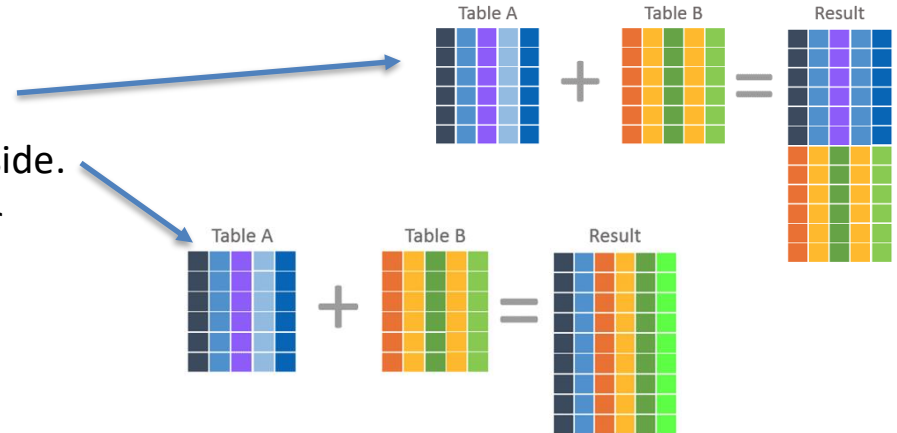
Consider Some Tables

Takes

ID	course_id	sec_id	semester	year	grade
00128	CS-101	1	Fall	2017	A
00128	CS-347	1	Fall	2017	A-
12345	CS-101	1	Fall	2017	C
12345	CS-190	2	Spring	2017	A
12345	CS-315	1	Spring	2018	A
12345	CS-347	1	Fall	2017	A
19991	HIS-351	1	Spring	2018	B
23121	FIN-201	1	Spring	2018	C+
44553	PHY-101	1	Fall	2017	B-
45678	CS-101	1	Fall	2017	F
45678	CS-101	1	Spring	2018	B+
45678	CS-319	1	Spring	2018	B
54321	CS-101	1	Fall	2017	A-
54321	CS-190	2	Spring	2017	B+
55739	MU-199	1	Spring	2018	A-
76543	CS-101	1	Fall	2017	A
76543	CS-319	2	Spring	2018	A
76653	EE-181	1	Spring	2017	C
98765	CS-101	1	Fall	2017	C-
98765	CS-315	1	Spring	2018	B
98988	BIO-101	1	Summer	2017	A
98988	BIO-301	1	Summer	2018	None

Student

ID	name	dept_name	tot_cred
00128	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
19991	Brandt	History	80
23121	Chavez	Finance	110
44553	Peltier	Physics	56
45678	Levy	Physics	46
54321	Williams	Comp. Sci.	54
55739	Sanchez	Music	38
70557	Snow	Physics	0
76543	Brown	Comp. Sci.	58
76653	Aoi	Elec. Eng.	60
98765	Bourikas	Elec. Eng.	98
98988	Tanaka	Biology	120

Columbia | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Consider a Subquery Tables

select *, (select name from student where student.id=takes.id) as name from takes;

Takes

ID	course_id	sec_id	semester	year	grade
00128	CS-101	1	Fall	2017	A
00128	CS-347	1	Fall	2017	A-
12345	CS-101	1	Fall	2017	C
12345	CS-190	2	Spring	2017	A
12345	CS-315	1	Spring	2018	A
12345	CS-347	1	Fall	2017	A
19991	HIS-351	1	Spring	2018	B
23121	FIN-201	1	Spring	2018	C+
44553	PHY-101	1	Fall	2017	B-
45678	CS-101	1	Fall	2017	F
45678	CS-101	1	Spring	2018	B+
45678	CS-319	1	Spring	2018	B
54321	CS-101	1	Fall	2017	A-
54321	CS-190	2	Spring	2017	B+
55739	MU-199	1	Spring	2018	A-
76543	CS-101	1	Fall	2017	A
76543	CS-319	2	Spring	2018	A
76653	EE-181	1	Spring	2017	C
98765	CS-101	1	Fall	2017	C-
98765	CS-315	1	Spring	2018	B
98988	BIO-101	1	Summer	2017	A
98988	BIO-301	1	Summer	2018	None

- Assume I wrote a function find_student_name(x)
  - Input is an x
  - Loops through all students and returns students with student.ID = x.
- The query with a subquery above is like:

result = []
For t in takes:
    new_r = t + find_student_name(t.id)
    result.append(new_r)

Student

ID	name	dept_name	tot_cred
00128	Zhang	Comp. Sci.	102
12345	Shankar	Comp. Sci.	32
19991	Brandt	History	80
23121	Chavez	Finance	110
44553	Peltier	Physics	56
45678	Levy	Physics	46
54321	Williams	Comp. Sci.	54
55739	Sanchez	Music	38
70557	Snow	Physics	0
76543	Brown	Comp. Sci.	58
76653	Aoi	Elec. Eng.	60
98765	Bourikas	Elec. Eng.	98
98988	Tanaka	Biology	120

Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.

- The nesting can be done in the following SQL query

  select $A_1, A_2, ..., A_n$
  from $r_1, r_2, ..., r_m$
  where P

  Note: Subquery MUST return
  - A single scalar if in the SELECT.
  - A Table if in the FROM.
  - If in the WHERE:
    - Either a scalar or a table.
    - Depending on the operation.

  as follows:

  - **From clause:** r_i can be replaced by any valid subquery

  - **Where clause:** P can be replaced with an expression of the form:

    B <operation> (subquery)

    B is an attribute and <operation> to be defined later.

  - **Select clause:**

    A_i can be replaced be a subquery that generates a single value.

Set Membership

Set Membership

- Find courses offered in Fall 2017 and in Spring 2018

      select distinct *course_id*
      from *section*
      where *semester* = 'Fall' **and** *year* = 2017 **and**
              course_id **in** (**select** *course_id*
                              from *section*
                              where *semester* = 'Spring' **and** *year* = 2018);

- Find courses offered in Fall 2017 but not in Spring 2018

      select distinct *course_id*
      from *section*
      where *semester* = 'Fall' **and** *year* = 2017 **and**
              course_id **not in** (**select** *course_id*
                              from *section*
                              where *semester* = 'Spring' **and** *year* = 2018);

Set Membership (Cont.)

- Name all instructors whose name is neither "Mozart" nor Einstein"

    select distinct *name*
    from *instructor*
    where *name* **not in** ('Mozart', 'Einstein')

- Find the total number of (distinct) students who have taken course sections taught by the instructor with *ID* 10101

    select count (**distinct** *ID*)
    from *takes*
    where (*course_id*, *sec_id*, *semester*, *year*) **in**
                        (**select** *course_id*, *sec_id*, *semester*, *year*
                         from *teaches*
                         where *teaches.ID*= 10101);

- Note: Above query can be written in a much simpler manner.
  The formulation above is simply to illustrate SQL features

Set Comparison

Set Comparison – "some" Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

  select distinct *T.name*
  from *instructor* **as** *T*, *instructor* **as** *S*
  where *T.salary* > *S.salary* **and** *S.dept name* = 'Biology';

- Same query using > **some** clause

  select *name*
  from *instructor*
  where *salary* > **some** (**select** *salary*
                                    from *instructor*
                                    where *dept name* = 'Biology');

Definition of "some" Clause

- F <comp> **some** $r \Leftrightarrow \exists\, t \in r$ such that (F <comp> t)
  Where <comp> can be: $<, \leq, >, =, \neq$

$$(5 < \textbf{some} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$$ (read: 5 < some tuple in the relation)

$$(5 < \textbf{some} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 = \textbf{some} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$$

$$(5 \neq \textbf{some} \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$$

(= **some**) $\equiv$ **in**
However, ($\neq$ **some**) $\not\equiv$ **not in**

Set Comparison – "all" Clause

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

> **select** *name*
> **from** *instructor*
> **where** *salary* > **all** (**select** *salary*
>   **from** *instructor*
>   **where** *dept name* = 'Biology');

Definition of "all" Clause

- F <comp> **all** $r \Leftrightarrow \forall\ t \in r$ (F <comp> *t)*

$(5 < \textbf{all}\ \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}\ ) = \text{false}$

$(5 < \textbf{all}\ \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}\ ) = \text{true}$

$(5 = \textbf{all}\ \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}\ ) = \text{false}$

$(5 \neq \textbf{all}\ \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}\ ) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$

$(\neq \textbf{all}) \equiv \textbf{not in}$
However, $(= \textbf{all}) \not\equiv \textbf{in}$

Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.

- **exists** $r \Leftrightarrow r \neq \emptyset$

- **not exists** $r \Leftrightarrow r = \emptyset$

Use of "exists" Clause

- Yet another way of specifying the query "Find all courses taught in both the Fall 2017 semester and in the Spring 2018 semester"

  select *course_id*
  from *section* **as** *S*
  where *semester* = 'Fall' **and** *year* = 2017 **and**
         exists (**select** *
              from *section* **as** *T*
              where *semester* = 'Spring' **and** *year* = 2018
                  and *S.course_id* = *T.course_id*);

- **Correlation name** – variable S in the outer query

- **Correlated subquery** – the inner query

Use of "not exists" Clause

- Find all students who have taken all courses offered in the Biology department.

    select distinct *S.ID*, *S.name*
    from *student* **as** *S*
    where not exists ((**select** *course_id*
                           from *course*
                           where *dept_name* = 'Biology')
                       except
                        (**select** *T.course_id*
                          from *takes* **as** *T*
                          where *S.ID* = *T.ID*));

- First nested query lists all courses offered in Biology
- Second nested query lists all courses a particular student took

- Note that $X - Y = \emptyset \iff X \subseteq Y$

- Note: Cannot write this query using = all and its variants

Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.

- The **unique** construct evaluates to "true" if a given subquery contains no duplicates .

- Find all courses that were offered at most once in 2017

  select *T.course_id*
  from *course* **as** *T*
  where unique (**select** *R.course_id*
                          from *section* **as** *R*
                          where *T.course_id= R.course_id*
                               and *R.year* = 2017);

Subqueries in the From Clause

Subqueries in the Form Clause

- SQL allows a subquery expression to be used in the **from** clause

- Find the average instructors' salaries of those departments where the average salary is greater than $42,000."

    select *dept_name*, *avg_salary*
    from (**select** *dept_name*, **avg** (*salary*) **as** *avg_salary*
             from *instructor*
             group by *dept_name*)
    where *avg_salary* > 42000;

- Note that we do not need to use the **having** clause

- Another way to write above query

    select *dept_name*, *avg_salary*
    from (**select** *dept_name*, **avg** (*salary*)
             from *instructor*
             group by *dept_name*)
             as *dept_avg* (*dept_name*, *avg_salary*)
    where *avg_salary* > 42000;

With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.

- Find all departments with the maximum budget

    with *max_budget* (*value*) **as**
        (**select max**(*budget*)
         from *department*)
    select *department.name*
    from *department*, *max_budget*
    where *department.budget = max_budget.value*;

Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
        with dept _total (dept_name, value) as
            (select dept_name, sum(salary)
             from instructor
             group by dept_name),
        dept_total_avg(value) as
            (select avg(value)
             from dept_total)
        select dept_name
        from dept_total, dept_total_avg
        where dept_total.value > dept_total_avg.value;
```

Scalar Subquery

- Scalar subquery is one which is used where a single value is expected

- List all departments along with the number of instructors in each department

select *dept_name*,
      (**select count**(*)
        from *instructor*
        where *department.dept_name* = *instructor.dept_name*)
      as *num_instructors*
from *department*;

- Runtime error if subquery returns more than one result tuple

Do Some Examples if not too Tired or Bored

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

Random Thoughts on Projects

COLUMBIA | ENGINEERING
The Fu Foundation School of Engineering and Applied Science

- The programming track will implement a simple, full stack web application.

**Full-stack Web Developer**

A full-stack web developer is a person who can develop both **client** and **server** software.

In addition to mastering HTML and CSS, he/she also knows how to:

- Program a **browser** (e.g. using JavaScript, jQuery, Angular, or Vue)
- Program a **server** (e.g. using PHP, ASP, Python, or Node)
- Program a **database** (e.g. using SQL, SQLite, or MongoDB)
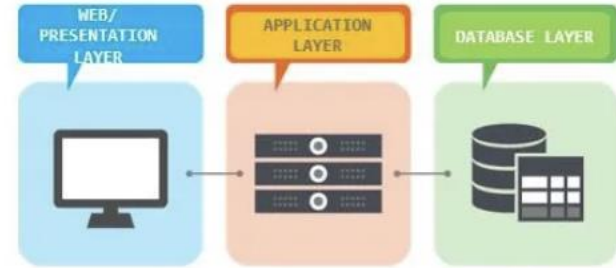
- The non-programming track will implement a data engineering and visualization process in a Jupyter notebook.

Columbia | Engineering
The Fu Foundation School of Engineering and Applied Science
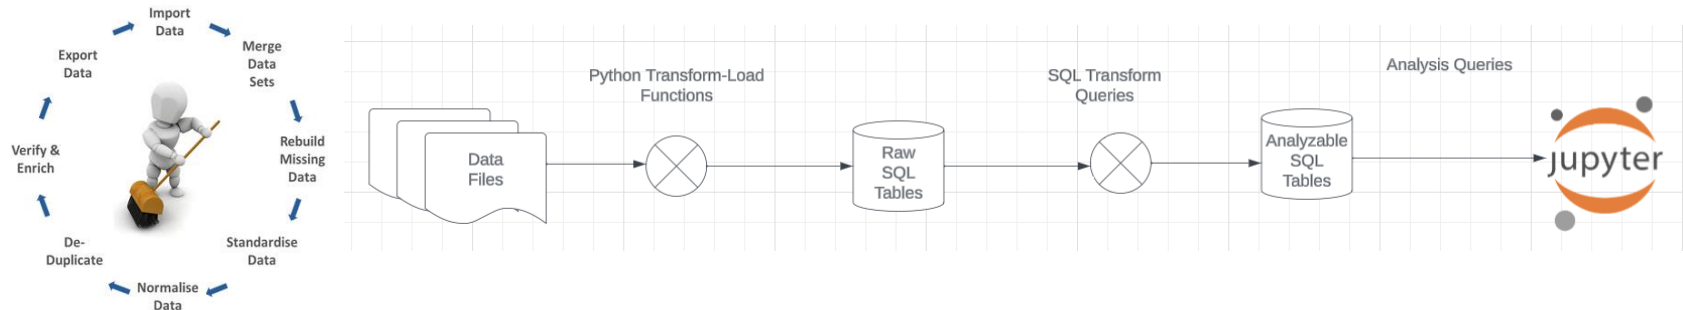
Approach

- You will find or generate an interesting dataset.

- I will provide templates and get started code for the

  - Full stack project

  - Data project

- Over the semester, starting with HW2, you will start to flesh out.

- A lot of my examples will use IMDB and Game of Thrones, but this data is huge.

- So, I used ChatCPT to generate some Lord of the Rings Data.

  - Scroll through ChatGPT history.

  - Show project-get-started.ipynb

- You can do something similar or find some interesting data.

Columbia │ Engineering
The Fu Foundation School of Engineering and Applied Science